

Scaling Strategy for Tinder-like Application

1. System Overview

A Tinder-like application must efficiently handle a **high read-heavy workload**, optimize for **low-latency user interactions**, and ensure **high availability and scalability**. Given the assumptions:

- **80 million MAU (Monthly Active Users)**
- **20 million DAU (Daily Active Users)**
- **100 million Peak DAU (5x DAU for high availability)**
- **QPS (API request estimation): 3,472 QPS**
- **Storage QPS: 80,000 QPS**

This document outlines a **FAANG-worthy scaling strategy** to ensure the application scales seamlessly under peak load conditions.

2. Database Scaling Strategy

Optimized Sharding Strategy

- **Geographical Partitioning:** Users are sharded based on their location (e.g., country/state/city). This ensures that users match with others within their geographical proximity, reducing cross-region data fetch latency.
- **User ID-Based Sharding:** Within a geographical region, further sharding is done using **User ID hashing** to distribute load evenly across database nodes.
- **Hot Users Handling:**
 - Popular profiles (celebrities, influencers) receive excessive matches/messages, leading to data skew.
 - Solution: Store hot users in a **separate high-performance database** and use a **consistent hashing technique** to distribute their profiles across multiple shards dynamically.
 - **Redis-backed caching layer** for frequently accessed hot profiles to minimize database reads.

3. Caching Strategy

Multi-Layered Caching

- **CDN (Content Delivery Network):** Used for serving **static content** like profile images, JavaScript, CSS.
- **Edge Caching:**
 - Implement **geo-distributed cache** using **CloudFront / Fastly**.

- Cache **frequent searches** and **match results** at the edge to reduce database hits.
- **Application-Level Caching (Redis / Memcached):**
 - **Profile Data:** Store **user profile** details in a Redis cache with **TTL-based expiration**.
 - **Match Lists:** Recent match lists are stored in cache to avoid repeated database queries.
 - **Swipe Actions:** Cache last **few hundred swipes** per user in Redis to reduce storage QPS.
- **Write-Through vs. Write-Back:**
 - **Write-Through Cache** for **profile updates**, ensuring consistency with the database.
 - **Write-Back Cache** for non-critical operations like user preferences.

4. Scaling for High Availability

Hybrid Approach: Replication + Sharding

- **Primary-Replica Replication:**
 - Read-heavy queries served via **read replicas** to handle high read QPS.
 - **Eventual consistency** is maintained via async replication to reduce write bottlenecks.
- **Auto-Scaling Database Clusters:**
 - DynamoDB/Cassandra for NoSQL-based user profiles & match preferences.
 - PostgreSQL/MySQL for transactional data like payments.

Load Balancing & Rate Limiting

- **Global Load Balancer (GLB):**
 - Routes traffic based on region to the nearest data center.
- **API Gateway with Rate Limiting:**
 - Limits API calls based on **subscription plans** (e.g., Tinder Gold users get more requests).

5. Handling Hot Users Problem

- **Separate Hot User Cluster:**
 - Profiles exceeding a threshold of requests per second are **offloaded to separate DB clusters**.
 - **Dedicated in-memory Redis cache** to store hot user profiles.
- **Adaptive Load Distribution:**
 - **AI-based predictive load balancing** to redirect traffic dynamically to less-loaded servers.
- **Bloom Filters:**

- Prevent redundant requests for the same hot user by storing **recent requests in a Bloom filter**.

6. Optimized Bandwidth & Resource Estimation

- **Outbound Data Estimation:**
 - Profile pictures (300KB avg per user, 7 images per user) + text data.
 - Match list responses, swipe actions, chat messages.
- **CPU & Memory Optimization:**
 - **CPU Cores** required based on request processing time estimation.
 - **Memory-based caching strategies** to reduce database pressure.